

A Comparison of Optimal Scanline Voxelization Algorithms

Tim Håkansson

Master of Science Thesis in Computer Science and Software Engineering

A Comparison of Optimal Scanline Voxelization Algorithms:

Tim Håkansson

LiTH-ISY-EX--20/5307--SE

Supervisor: **Harald Nautsch**

ISY, Linköping University

Åsa Detterfelt

MindRoad

Jens Ogniewski

MindRoad

Examiner: **Ingemar Ragnemalm**

ISY, Linköping University

*Division of Information Coding
Department of Electrical Engineering
Linköping University
SE-581 83 Linköping, Sweden*

Copyright © 2020 Tim Håkansson

Abstract

This thesis presents a comparison between different algorithms for optimal scanline voxelization of 3D models. As the optimal scanline relies on line voxelization, three such algorithms were evaluated. These were *Real Line Voxelization* (RLV), *Integer Line Voxelization* (ILV) and a 3D Bresenham line drawing algorithm. RLV and ILV were both based on voxel traversal by Amanatides and Woo. The algorithms were evaluated based on runtime and the approximation error of the integer versions, ILV and Bresenham. The result was that RLV performed better in every case, with ILV being 20-250% slower and Bresenham being 20-500% slower. The error metric used was the Jaccard distance and generally started at 20% and grew up towards 25% for higher voxel resolutions. This was true for both ILV and Bresenham. The conclusion was that there is no reason to use any of the integer versions over RLV. As they both performed and approximated the original 3D model worse.

Acknowledgments

I would like to thank my supervisors at Mindroad, Åsa Detterfelt and Jens Ogniewski, for guiding me through this thesis. I would also like to thank my colleagues at Mindroad for the great company during our breaks. A thanks also goes out to my examiner Ingemar Ragnemalm, both for the help provided during the thesis, but also for the great courses he provided during my studies. A special thanks goes out to my family for all their love and support during my education.

Linköping, June 2020
Tim Håkansson

Contents

1	Introduction	1
1.1	Background	1
1.2	Aim	2
1.3	Research Questions	2
1.4	Delimitations	2
1.5	Mindroad	3
2	Background	5
2.1	GPU-Programming	5
2.1.1	OpenGL	6
2.1.2	CUDA	6
2.2	Amazon Web Service	7
2.3	Voxelization	7
2.4	Raycasting	8
3	Theory	9
3.1	Line Voxelization	9
3.1.1	Real Line Voxelization	10
3.1.2	Integer Line Voxelization	11
3.1.3	3D Bresenham Algorithm	11
3.2	Model Voxelization	12
3.2.1	Floating-Point Optimal Scanline	12
3.2.2	Integer Optimal Scanline	14
3.2.3	Rasterization	14
3.2.4	Depth Buffer	15
3.3	Error Analysis	15
3.3.1	Relative Error	15
3.3.2	Jaccard Distance	16
4	Method	17
4.1	Implementation	17
4.1.1	CUDA-OpenGL Interoperability	17
4.1.2	Floating-Point Voxelization	18

4.1.3	Integer Voxelization	19
4.1.4	Bresenham Algorithm	20
4.1.5	Voxel Rendering	21
4.2	Evaluation	22
4.2.1	Performance Analysis	22
4.2.2	Error Analysis	23
5	Results	25
5.1	Voxelization	25
5.2	Performance Analysis	25
5.3	Error Analysis	30
6	Discussion	31
6.1	Results	31
6.1.1	Performance of Models	31
6.1.2	Performance Between RLV and ILV	32
6.1.3	Performance Between ILV and Bresenham	32
6.1.4	Performance Data	32
6.1.5	Error Analysis	33
6.2	Method	34
6.2.1	Performance Analysis	34
6.2.2	Error Analysis	34
6.3	Future Work	34
7	Conclusion	35
7.1	Research Questions	35
7.2	Choice of Algorithm	36
A	Appendix	39
A.1	3D Bresenham Algorithm	39
A.2	6-Connected Bresenham Algorithm Modification	40
A.3	CUDA-OpenGL Interoperability	41
A.4	Voxelizations	42
A.5	Performance Data	45
A.6	Voxelization Error	45
Bibliography		51

1

Introduction

This chapter gives the reader an introduction to voxels, its uses and ways to create voxels from a model. It also serves the purpose of motivating the thesis, as well as describing the problem it is out to solve. Finally, it defines the delimitations of the thesis and presents the company, at which this thesis was conducted at.

1.1 Background

In computer graphics, a 3D model is built up of triangles which describes the surface of the model. This means it does not store any information to represent the inside of the model. In order to represent the inside, a volumetric data structure is needed. Usually, this is done by storing *voxels* in a uniform 3D grid. A voxel can be described as data at a position in a grid. This can be any form of data, such as occupancy, color, material or density.

As voxels are just a way to represent volumetric data, they have been used in a wide range of applications, such as global illumination [2], medical imaging [3, 4] and collision culling [5]. These applications either store a voxel representation of models (especially true for medical imaging) or need to convert models into voxels. This process is called *voxelization* and has been widely studied in the past. Some methods for voxelizations include triangle-box intersection [6], rasterization [7] and depth buffer [20].

Recently, a study was published by Zhang et al. [9] proposing a new method to voxelize a model. The basics of the method is to voxelize the model by performing line voxelization at different stages. This method is the basis of this thesis and will further be called the *optimal scanline*.

1.2 Aim

This thesis aims to do an investigation of how the line voxelization algorithm affects the performance of the optimal scanline. This includes both floating-point and integer line voxelizations.

The thesis also investigates the approximation error caused by using the integer versions of the algorithm. This was done in [9], but the authors only compared the voxel count. Which means if a voxel is moved somewhere else, it would not produce any error. As such, this thesis aims to measure this error in another metric, which can describe those errors.

1.3 Research Questions

With the aim of the thesis defined, two research questions are formed as a baseline of the thesis. These are presented below:

1. Which line voxelization algorithm performs best for the optimal scanline?
2. How great is the approximation error of the integer versions of the optimal scanline?

A better performance in this case is defined as how fast the execution of the voxelization is. It is not a metric of the error or the memory usage of the algorithms.

1.4 Delimitations

The implementation of the thesis ran on *Amazon Web Services* (AWS), on a computer running Ubuntu 18.04 with an NVIDIA Tesla T4 graphics card with driver version 440.59.

The rendering of the voxelization was done using OpenGL. As the focus was not on supporting older devices, the OpenGL version 4.6 was used.

The computing language of choice was CUDA, as such, other languages were not considered for the implementation. Again, as there is no need to support older devices, CUDA 10.2 was used.

The line algorithms that were evaluated were limited to a voxel traversal algorithm by Amanatides and Woo [1], its integer version and Bresenham. As Bresenham is predominantly a 2D line drawing algorithm, it was extended to 3D using a method proposed by Liu and Cheng [10].

The evaluations were also limited to three different models and voxel grid resolutions between 128-2048.

1.5 Mindroad

This master thesis was conducted on behalf of MindRoad. MindRoad is a software company which specializes in embedded systems, web development and mobile applications. They also provide courses in software development, such as C++, Python, GPU-programming, Linux and driver development.

2

Background

This chapter serves the purpose of giving the reader a better understanding of central concepts throughout the thesis. It will give an introduction to GPU-programming, AWS, voxelization and raycasting.

2.1 GPU-Programming

In GPU-programming, there are two different subjects to discuss, rendering and computing. Rendering is where the GPU performs calculations, which results in some visual output. This most often means taking triangles, performing different transformations on them and then drawing them to a screen. Computing is similar, but instead of resulting in graphics on a screen, the results are stored in video memory as data. This is useful when algorithms can make use of the GPU's parallelization.

In order for the GPU to do any form of calculations, an API is needed to interact with it. For rendering, there are API:s such as OpenGL, DirectX and Vulkan. API:s used for computing include CUDA and OpenCL. All the rendering API:s above also support computing, but it is not what their general use is.

The chosen API:s for this thesis were OpenGL and CUDA, and are presented in the next sections.

2.1.1 OpenGL

OpenGL [11], short for Open Graphics Library, is technically not an API but an open specification which defines functions used for 2D and 3D graphics. As it is a specification, there are implementations for most platforms, including Windows, Linux, MacOS¹, Android and iOS¹.

To render with OpenGL, shaders are used to transform and color primitives, such as triangles. The shaders are, in a sense, programs which are executed on the GPU. When rendering a model, the GPU can run these shaders in parallel either per vertex, triangle or pixel, depending on the shader. These are coded in GLSL, which supplies the developer with a wide range of vector math functions.

OpenGL requires two shader programs in order to operate, the vertex and fragment shader. The input to the vertex shader is, of course, vertices. Often these vertices contain a position, normal and texture coordinate, and define a corner of a triangle. The vertex shader then performs transformations on the vertices in order to move or in any way alter the triangles. The result is then sent to a rasterizer, which turns these triangles into fragments. These fragments are then the input of the fragment shader. In simplest terms, fragments are pixels on the screen before they have been processed by the fragment shader. The fragment shader is then able to perform calculations which changes the color of the fragment. This can be based on normals, light positions, distance to the camera, etc. The resulting color is then the color of the pixel on the screen.

There are two more programmable shaders in OpenGL, tessellation and geometry. These are not required to render to the screen and will therefore not be discussed further in this thesis.

2.1.2 CUDA

CUDA is an API for performing general computing on a GPU. This API is being developed by NVIDIA and is only available on NVIDIA GPU:s. However, it is cross-platform in terms of operating systems, as it is supported on Windows, Linux and MacOS.

CUDA is designed as a C/C++ extension and compiles using *NVIDIA CUDA Compiler* (NVCC). This compiler in turn makes use of a host compiler (such as g++) in order to compile the C/C++ code [14]. As such, there should not be any performance difference in the CPU code by using NVCC.

CUDA also has two different API:s which the developer can interact with, the driver and the runtime API. The difference between them is minor, runtime CUDA is higher-level and is able to link the CUDA kernels into the compiled executable. Kernels are basically CUDA programs, much like shaders are OpenGL programs. The driver API on the other hand is lower-level and requires the program to read the CUDA kernels from an externally compiled cuda binary.

¹OpenGL is currently deprecated (but still works) on Apple devices in favor of Metal [12, 13]

2.2 Amazon Web Service

AWS [15] is a set of cloud computing services provided by Amazon. These services can provide the user with GPU computing, database management, cloud storage and many more services.

To interact with these services the protocol SSh is used. This provides a way of connecting to a server with a terminal interface. A problem with this, for the thesis, is that rendering is required to test the voxelization results. In order to get a graphical interface on an AWS, Thinlinc [16] was used. ThinLinc is a software which runs a desktop environment on a server and transfers the display over a network. However, ThinLinc in itself does not support 3D hardware acceleration, which is needed by OpenGL. As such, VirtualGL [17] was used when running the voxelization over ThinLinc. This enables ThinLinc to run OpenGL with 3D hardware acceleration. To run a program with VirtualGL simply prepend the program with `vglrun`. For example, run `vglrun glxgears` to run the Linux GLX demo in ThinLinc.

2.3 Voxelization

In 2D, the process of turning triangles and other geometrical shapes into pixels is called *rasterization*. Rasterization is required to render shapes onto a screen, as monitors cannot handle continuous shapes and have to discretize them into pixels. An example of rasterization can be seen in Figure 2.1. Voxelization is the 3D version of rasterization, where instead of marking pixels in a 2D plane, it marks voxels in a 3D space. Recall that a voxel can store any form of data, be it color, material or density. What is stored is dependent on the use case of the voxels. In this thesis, only occupancy is stored, that is if the voxel exists or not.




Figure 2.1: Rasterization of a triangle

With voxelization defined there are two other topics that need discussing. The first of which is if the algorithm supports surface or solid voxelization. *Surface voxelization* means the algorithm only voxelizes the outer surface of the model, and therefore leaves the voxelization with an empty interior. *Solid voxelization* is the opposite of that, filling the interior with voxels. In this thesis, only surface voxelization will be considered, due to the optimal scanline only supporting it.

The other topic is the *connectivity* of the voxelization. There are three types of connectivity in a 3D voxelization, 6-, 18- and 26-connected. The number signifies how many possibilities there are for two voxels to be considered connected. A 2D example is seen in Figure 2.2, where it is either 4- or 8-connected. Similarly in 3D, two voxels are 26-connected if they share either a face, an edge or a corner between them. Two voxels are 18-connected if they share a face or an edge between them. Finally, two voxels are 6-connected if they share a face between them. The authors of the optimal scanline claim the method can produce any voxelization connectivity, but mainly focuses on 6-connected. However, for the scanlines to work, 6-connected line voxelization is required.



Figure 2.2: Connectivity in the 2D case, where the light gray squares are connected to the dark gray square. The left figure is 4-connected, whilst the right is 8-connected.

2.4 Raycasting

Raycasting is a rendering technique where lines, often called rays, are being sent out from the camera. If these rays intersect with an object in the world, that part of the object should be visible on the screen. Information such as color, depth and material of the object can be retrieved by looking at where the ray intersected.

An early example of raycasting is its use in Wolfenstein 3D [18]. In this game, a ray is being sent out for each column of the screen. If a ray hits a wall, the game renders a vertical line with the height depending on the distance to the wall. The color of the pixels in the line is then dependent on the texture of the wall.

Nowadays, raycasting can be performed for each pixel on the screen. That is, for each pixel, send out a ray from the camera and color the pixel depending on what its ray intersects with. As such, there is no need to perform any mesh rendering and only objects that are visible are being rendered. For some use cases, this might be an optimal solution. It is however an expensive operation if the use case is to render highly detailed triangle meshes, since it is difficult to find which triangle the ray intersects with. This is because there are potentially thousands of triangles per model and millions of rays to test. Cases where raycasting is better to use are voxel grids, as there are algorithms to traverse it without evaluating every voxel. One such technique is called *raymarching* which marches through all voxels the ray goes through. This is then terminated when a voxel is found or when the ray exits the scene.

3

Theory

This chapter presents the literature used in order to answer the research questions. It starts off with defining line voxelization and the three algorithms used. Then the optimal scanline is presented together with other voxelization techniques. Finally, some error metrics are presented, one of which is used as part of the final results.

3.1 Line Voxelization

Line voxelization is a way of generating voxels based on a 3D line. This is needed for the optimal scanline method to work, as this is how the scanlines are generated. In this thesis, line voxelization defines all voxels in a grid being touched by a line. This is similar to how lines are drawn to a screen in 2D. An example of line voxelization can be seen in Figure 3.1. To determine which voxels are touching the line, three algorithms are introduced in the upcoming sections.




Figure 3.1: Line voxelization, where all cells touching the line is marked as voxel. This works the same in 3D, but with an additional axis.

3.1.1 Real Line Voxelization

Amanatides and Woo [1] proposed a method of ray marching voxels in a uniform grid. This can also be used to voxelize lines and will further be called *Real Line Voxelization* (RLV). The basis of the algorithm is the line equation

$$p = p_0 + vt,$$

where p is a position on the line, p_0 is the start position of the line, v is the direction of the line (normalized) and t is how much the direction is scaled.

When the algorithm starts, it initializes t as

$$t_x = \frac{p_{1,x} - p_{0,x}}{v_x},$$

where $p_{1,x}$ is the x-position of the next voxel in the x-axis, as seen in Figure 3.2. This is calculated similarly for each component of t . Each iteration of the algorithm evaluates t_{min} as $\min(t_x, t_y, t_z)$. Let's assume $t_{min} = t_x$. This means the next voxel is in the x-direction. Again, from Figure 3.2, $p_{1,x}$ is closer to p_0 compared to $p_{1,y}$ and therefore has a smaller t value. This means the next voxel would be placed to the right of p_0 . The t for the next iteration is calculated by subtracting all of its components by t_{min} .

Finally assign

$$t_x = \frac{|p_{i,x} - p_{i-1,x}|}{v_x} = \frac{1}{v_x},$$

where $p_{i,x}$ is the x-position of the i 'th voxel in the x-direction. Here the grid is assumed to have a voxel size of 1, meaning $|p_{i,x} - p_{i-1,x}| = 1$. After this the iteration is restarted. The iteration is then terminated when the sum of all t_{min} exceeds the length of the line.




Figure 3.2: How $p_{1,x}$ and $p_{1,y}$ are defined given a start position, p_0 , and a direction, v

3.1.2 Integer Line Voxelization

Integer Line Voxelization (ILV) follows the same structure as RLV except it avoids the floating-point arithmetics and divisions. The changes needed to avoid floating-points are adapted from Zhang et al. [9].

Having RLV as a basis, ILV requires three changes. Firstly, the initial t is calculated from the center of the start voxel, which means

$$p_{1,x} - p_{0,x} = \frac{1}{2},$$

and therefore

$$\begin{cases} t_x = \frac{1}{2\Delta X} \\ t_y = \frac{1}{2\Delta Y} \\ t_z = \frac{1}{2\Delta Z} \end{cases}$$

where ΔX denotes the length of the line in the x-axis. ΔY and ΔZ are defined similarly.

Secondly, since only the relative sizes of t 's components are needed, the equation can therefore be multiplied with

$$2\Delta X\Delta Y\Delta Z,$$

to avoid fractions. As such, the integer version of t_x is denoted T_x and can be described as

$$T_x = \Delta Y\Delta Z. \quad (3.1)$$

T_y and T_z are described similarly.

Lastly, the iteration of the algorithm follows RLV, replacing t with T , but the assignment to t_x in the end is instead

$$T_x = 2\Delta Y\Delta Z,$$

if $T_{min} = T_x$. This follows Equation 3.1, but multiplied with 2. The reason is that before T was calculated to traverse half a voxel, multiplying it by 2 then traverses a whole voxel.

3.1.3 3D Bresenham Algorithm

The 3D version of the *Bresenham* algorithm presented by Liu and Cheng [10] follows the original by Bresenham [19] rather nicely. The 3D algorithm starts off with a few assumptions about the line, $\Delta X \geq \Delta Y \geq 0$ and $\Delta X \geq \Delta Z \geq 0$, these are defined the same as in the previous section.

In the original Bresenham an error in the y-axis is initialized to

$$e_y = 2\Delta Y - \Delta X.$$

Then for each iteration, if the current x is greater than X_1 , terminate the iteration. Otherwise set the current voxel and increment x by one. Following that, check if the error of the line is greater than 0. If it is, increase y by one and decrease the error, otherwise increase the error. Then restart the loop. The change in the error is defined as

$$\begin{cases} e'_y = e_y + 2(\Delta Y - \Delta X) & , e_y \geq 0 \\ e'_y = e_y + 2\Delta Y & , e_y < 0 \end{cases}.$$

To extend the algorithm to 3D, Liu and Cheng added another error, e_z , which kept track of when z should increase. This works the same way as e_y and is evaluated after e_y . A pseudo code of the algorithm can be found in Appendix A.1.

3.2 Model Voxelization

As mentioned in the introduction, voxelization is a way of turning a 3D model or scene into voxel data. This can be done in a wide range of ways and is still being researched today. Four such techniques, two of which are based on the optimal scanline, are presented in the following sections.

3.2.1 Floating-Point Optimal Scanline

One of the more recent works in voxelization include work done by Zhang et al. [9]. This algorithm, called optimal scanline, uses line voxelization as its core concept. All its calculations are done for each triangle of the model. The algorithm first sorts the vertices of a triangle based on its most dominant axis. This axis is defined as the x-, y- or z-axis which most aligns with the triangle's normal. It can be determined by choosing the axis with the greatest absolute value of the normal's components. That is, choose the axis which matches $\max(|n_x|, |n_y|, |n_z|)$, where n is the normal of the triangle. The dominant axis will further be assumed to be the z-axis.

With the vertices of the triangle sorted, it performs line voxelizations between each of the vertices using RLV. This results in all the edges of the triangle being voxelized.

In order to fill the interior of the triangle, it splits the edges into slices in the z-axis. Where each edge has the same integer z-value. Then, for each slice, perform 2D line voxelization between the edges of the triangle. This is shown in Figure 3.3. The figure also shows that not every edge voxel needs to contain a scanline endpoint. That is where the optimal keyword of the algorithm comes in. The authors propose a theorem stating that, there exists a distance, l , where there cannot exist a voxel between two parallel lines. This distance can be seen in Figure 3.4. Intuitively, the distance can be seen to be the length of the voxel's diagonal projected onto the scanline direction, d_{sl} . Furthermore, the length, l , can be calculated as $|d_{sl,x}| + |d_{sl,y}|$. To prove this the scanline can be classified into four cases, all the combination of the signs of d_{sl} . The result is the four following




Figure 3.3: Left figure shows a slice of the edges of the triangle together with its voxelization shown in gray. The right figure shows how the optimal scanlines are chosen and the resulting voxelization in light gray.




Figure 3.4: A voxel, in gray, being intersected with two parallel lines with a distance l between them. The direction of the scanline is labeled d_{sl} .

equations:

$$\begin{cases} l = ((0, 1) - (1, 0)) \cdot d_{sl} & , d_{sl,x} < 0 < d_{sl,y} \\ l = ((1, 0) - (0, 1)) \cdot d_{sl} & , d_{sl,y} < 0 < d_{sl,x} \\ l = ((0, 0) - (1, 1)) \cdot d_{sl} & , d_{sl,x}, d_{sl,y} < 0 \\ l = ((1, 1) - (0, 0)) \cdot d_{sl} & , 0 < d_{sl,x}, d_{sl,y} \end{cases}$$

The values $(0, 0)$, $(1, 1)$, $(1, 0)$ and $(0, 1)$ are the corners of the voxel. They are used to calculate the diagonal that aligns with d_{sl} . Calculating the dot product of each of the equations results in $l = |d_{sl,x}| + |d_{sl,y}|$.

3.2.2 Integer Optimal Scanline

The way integer voxelization of the optimal scanline is handled is a bit different from the floating-point one. Since the scanline direction and the scanline length are both floating-points, these cannot be used. Instead, it makes use of an iterative approach, where it steps through the edge voxels of the triangle until the voxel is too far away from the previous scanline. How exactly it was derived can be found in [9], but the final theory is presented below (assuming z is the most dominant axis).

To start off, two boundary variables are defined, called C_{lower} and C_{upper} . These are defined as

$$\begin{cases} C_{lower} = \Delta Y X^a - \Delta X Y^a - |\Delta X| - |\Delta Y| \\ C_{upper} = \Delta Y X^a - \Delta X Y^a + |\Delta X| + |\Delta Y| \end{cases},$$

where (X^a, Y^a) and (X^b, Y^b) are the endpoints of the previous scanline and $\Delta X = X^b - X^a$, $\Delta Y = Y^b - Y^a$. Each iteration, the algorithm steps the edge voxelization by one voxel and calculate a variable called C_k to be

$$C_k = \Delta Y X_k - \Delta X Y_k,$$

where (X_k, Y_k) is the k'th voxel on the edge, with $k=0$ being the previous scanline endpoint. If this variable is outside the boundaries, C_{lower} and C_{upper} , the voxel is too far away from the previous scanline and (X_{k-1}, Y_{k-1}) is chosen as the next scanline endpoint. This iteration is performed separately for both the edges. As a triangle has three edges and not two, the voxels on the edge from v_1 to v_2 are merged with the voxels on the edge from v_2 to v_3 . Those two edges are therefore counted as a single edge.

3.2.3 Rasterization

Another technique used for surface voxelization utilizes the GPU rasterizer in order to optimize the voxel generation. One such technique was presented by Crassin and Green [7]. The basic idea is that, for every triangle in the model, find the triangle's most dominant axis and render it from that direction. In practice, this means swapping the dominant axis with the z-axis in a vertex shader. The triangle is then sent to the rasterizer which outputs fragments of the triangle. The voxel coordinate is then simply the position of the fragment and its depth value. Then the axes are swapped back from before in order to place it correctly in the scene. Finally it writes the coordinate to a 3D texture which stores the voxel data. Note that all rasterization is done in a framebuffer with the same width and height as the voxel grid resolution. There are however problems with the algorithm creating holes in the voxelization in some instances. It can be resolved by using a technique called conservative rasterization. This involves marking all fragments which are touched by the triangle.

3.2.4 Depth Buffer

One way of performing solid voxelization is to make use of the depth buffer [20]. This is done by rendering the entire model from six different directions, positive x, y and z, and negative x, y and z. In these rendering steps, only the depth buffer is needed. Then the voxelization is defined as all voxels which are within all the depth buffers.

One way to do this, is to iterate through the entire grid and marking all voxels which are within the buffers as occupied. This is however rather slow for higher resolutions, as the complexity grows cubically. Another way would be to choose one axis, let's say the z-axis, and iterate through its buffer x- and y-coordinates. Then for each coordinate, there is a minimum and maximum z-position which is determined by the two depth buffers in the z-axis. This means the algorithm only needs to iterate through these two values instead of all the z-values. As most models do not fill the entire voxel grid, this reduces the runtime of the algorithm in the average case.

One obvious problem of the algorithm is that it needs to render the object six times to voxelize it. For convex shapes, this can be reduced to two, as the shape can be fully described by rendering it from the front and back.

3.3 Error Analysis

Calculating the approximation errors can be done using several methods. The choice of method depends on the use case. This thesis presents two different methods to analyze the error. These are relative error, the one used in [9], and the Jaccard distance.

3.3.1 Relative Error

The *relative error* can be described as

$$e_{re} = \frac{|x - y|}{|x|} = \left| \frac{x - y}{x} \right|,$$

where x is the actual value and y is the approximation. This was the method used by Zhang et al. [9], where they set x to the total voxel count of the floating-point version and y to the total voxel count of the integer version. A problem with this is that if a voxel is moved to another location, the error would be considered 0. This can be seen in Figure 3.5, where the relative error of the voxel count is 0. Therefore, a different error metric is required in order to describe this error.




Figure 3.5: Difference between RLV and ILV. White voxels are only in RLV, gray voxels are only in ILV and the dark gray voxels are in both. The voxel count is the same for both.

3.3.2 Jaccard Distance

The *Jaccard distance* is an error metric which aims to measure differences in sets. It is derived from the Jaccard similarity presented in [21], where it instead measures how similar two sets are. The Jaccard similarity is defined as the size of the intersection divided by the size of the union of the two sets. This results in a value between zero and one, which can be interpreted as a percentage of how similar two sets are. The Jaccard distance is defined as one subtracted by the Jaccard similarity. It can also be defined as the symmetric difference divided by the union of the two sets. Both of which are equivalent.

4

Method

This chapter describes the implementation and application of the theory in more detail. As such, it will build upon the literature in order to more precisely describe how the project was implemented. The chapter will also give a description of how the evaluation of the different algorithms was performed.

4.1 Implementation

The full source code of the project can be found at [22]. The implementation was done in C++ with OpenGL and CUDA. CUDA handled the voxelization and OpenGL was used to render the results of it. The project was implemented in Linux, but can of course be ported to other platforms. The implementation details which were not covered in the literature will be explained in the following sections.

4.1.1 CUDA-OpenGL Interoperability

During development, some form of visual feedback was required in order to verify and debug the results of the voxelization. As such, a CUDA-OpenGL interoperability had to be implemented. This meant creating a 3D texture in OpenGL and binding it to CUDA in order to write to it. A simplified source code of this can be seen in Appendix A.3, which is adapted from [23]. This method links the memory of the 3D texture to a CUDA array, meaning no data is duplicated. This is required, as texture sizes of up to 8 GB were used.

4.1.2 Floating-Point Voxelization

The first implementation detail that needs to be explained is how the scanline direction, d_{sl} , as shown in Figure 3.4, was calculated. Let's call the three vertices of a triangle v_1 , v_2 and v_3 , which are sorted in the most dominant axis (assumed to be the z-axis). The scanline direction was calculated in two different ways, depending on how the vertices were positioned. First, if all vertices were in the same z-slice, any direction would suffice. In the implementation, the edge between v_1 and v_3 was chosen. If the vertices were not in the same z-slice, the gradient of the triangle with respect to z was used. This was calculated by first solving the plane equation for z :

$$D = n_x x + n_y y + n_z z \rightarrow z = \frac{D - n_x x - n_y y}{n_z},$$

where n is the normal of the triangle and D describes the position of the plane. The gradient was then the partial derivatives of the equation with respect to x and y :

$$d = \left(\frac{-n_x}{n_z}, \frac{-n_y}{n_z} \right)$$

Which was the resulting scanline direction.

In both cases, the direction also needed to be normalized, meaning the final direction was

$$d_{sl} = \frac{d}{|d|}.$$

Each iteration of the algorithm, two scanline endpoints were calculated by reverse projecting the scanline direction onto the triangle's edges. The scanline direction was scaled by l_i , which was increased by l each iteration. Recall from the theory that $l = |d_{sl,x}| + |d_{sl,y}|$. Using the reverse projection, the endpoints were calculated as

$$v_p = v_1 + d_e \frac{l_i}{d_e \cdot d_{sl}}, \quad (4.1)$$

where d_e is the normalized direction of the edge it projects to. Note that this only works with the edges from v_1 to v_2 and v_1 to v_3 . So a special case was needed for the edge from v_2 to v_3 . This was solved by recalculating v_2 such that

$$\begin{cases} 0 = (v'_2 - v_1) \cdot d_{sl} \\ v'_2 - v_2 = (v_3 - v_2)t \end{cases}.$$

That is, $v'_2 - v_1$ is perpendicular to d_{sl} and $v'_2 - v_2$ is parallel with $v_3 - v_2$. An example of this can be seen in Figure 4.1. Solving the equation resulted in

$$v'_2 = v_2 + (v_3 - v_2) \frac{(v_1 - v_2) \cdot d_{sl}}{(v_3 - v_2) \cdot d_{sl}}.$$

This was used instead of v_1 in Equation 4.1, when projecting to the edge between v_2 and v_3 .




Figure 4.1: v'_2 is calculated by finding the intersection between the edge's line and the perpendicular line of d_{sl}

4.1.3 Integer Voxelization

In the integer version there were two edge cases which created holes in the voxelization. In order to solve these problems first recall the values C_{lower} , C_{upper} and C_k from the theory in Section 3.2.2.

The first problem occurred when a new slice was started and an example is shown to the left in Figure 4.2. More specifically, it occurred when the next voxel was behind the current scanline, meaning it should have been included in it. This was resolved by checking if the value of C_1 was on the other side of C_0 relative to C_{end} . That is, if $C_0 < C_1 < C_{end}$ or $C_{end} < C_1 < C_0$. Here C_{end} is defined as the C_k of the last voxel on the triangle's edge. To test if the two values were on different sides, following equation was evaluated:

$$(C_1 - C_0) * (C_{end} - C_0) \leq 0.$$

If this condition was met, either C_{upper} or C_{lower} was set to C_0 depending on if C_{end} was greater or less than C_0 respectively. This was done in the same step as when the next endpoint for the scanline is calculated.

The other case where holes occurred was when (X_1, Y_1) for both edges were on different sides of the scanline. This sometimes happened for triangles which had a very acute angle at v_1 . An example of this can be seen to the right in Figure 4.2. This problem was similarly detected as before by

$$(C_1^a - C_0) * (C_1^b - C_0) \leq 0,$$

where C_1^a and C_1^b are C_1 for the two edges. When this was the case, ΔX and ΔY were set similarly to how the scanline direction was chosen for the floating-point version. That is, the gradient of the triangle was calculated, but the division by n_z was removed to avoid floating-point operations. However, this resulted in a scanline direction and not a difference between scanline endpoints. So it was also rotated by 90° , as the scanline should always be perpendicular to the scanline direction. The result was $\Delta X = N_y$ and $\Delta Y = -N_x$, where N is the unnormalized normal of the triangle. These values were then used when determining the next scanline endpoints.




Figure 4.2: Voxel being missed by the scanlines. The dark gray rectangles are the edges of the triangle and the light gray rectangles are the voxels of the scanlines.

4.1.4 Bresenham Algorithm

Worth noting is that Bresenham only works when ΔX is positive and greater than ΔY and ΔZ . To allow for negative directions, three changes were required. First, whenever x , y or z increased by one, they were instead decreased by one if the difference was negative. Another change was to set ΔX , ΔY and ΔZ to their absolute value. Finally, the iteration was terminated when the x -position of the voxel equals X_1 , where X_1 is the last x -position on the line. This would however miss the last voxel on the line, which was resolved by setting the last voxel after the loop.

To solve the problem where ΔX has to be greater than the other differences, a variable which keeps track of this axis was introduced. This variable will be called A and was initialized to

$$\begin{cases} A = 0 & , \max(\Delta X, \Delta Y, \Delta Z) = \Delta X \\ A = 1 & , \max(\Delta X, \Delta Y, \Delta Z) = \Delta Y \\ A = 2 & , \max(\Delta X, \Delta Y, \Delta Z) = \Delta Z \end{cases}$$

Then whenever the ΔX or x was needed, instead the variable was indexed using A . For example $\Delta[A]$, would give the greatest difference. The indices of the other two axes was calculated by increasing A by one and two respectively, and then taking the modulus 3 of them.

Another thing missing from the theory is that Bresenham is not 6-connected. This can be seen in Figure 4.4. As the optimal scanline requires 6-connected lines in order to fill the interior, the Bresenham algorithm had to be modified to support this. The first intuitive way to do this would be to add a voxel at (x, y, z) whenever y or z increases in the algorithm. This however, caused the voxelization to not follow the line correctly, as seen to the left in Figure 4.4. This was solved by checking if the error in y was greater than ΔY , then instead of voxelizing the right voxel, voxelize the top voxel. The result can be seen to the right in Figure 4.4. The same check was performed on the z -axis. This modification of the algorithm in Appendix A.1 can be seen in Appendix A.2.

The final problem with the Bresenham algorithm was that the integer version of the scanline algorithm required being able to get the next voxel in the line. The

problem was that the 6-connected Bresenham could generate up to three voxels each iteration. One way to solve the problem would be to run one iteration of Bresenham and store all voxels generated in that iteration in a list. Then the next time a new voxel is needed, return an unused voxel from the list. The way it was solved for this thesis however, was to keep track of where the last voxel was returned in the iteration. Then the next time a voxel was needed, the iteration resumed where it left off last.




Figure 4.3: Original Bresenham algorithm which does not voxelize all voxels touching the line




Figure 4.4: Wrong and correct 6-connected Bresenham algorithm to the left and right respectively. The light gray rectangles shows the added voxels not in the original Bresenham algorithm.

4.1.5 Voxel Rendering

In order to test the results of the optimal scanline, some form of rendering was needed. This thesis only considers rendering which results in cubes in a uniform grid.

One approach would be to iterate through the 3D texture on the CPU and render a cube for every existing voxel. This would be bad for numerous reasons. First off, the data created on the GPU would need to be copied over to the CPU, which creates significant overhead. Especially since the data could get up towards 8 GB in size for the highest resolution. Secondly, rendering all voxels (even occluded ones) would slow down the rendering significantly.

Another way would be to write each voxel coordinate to a list on the GPU when voxelizing the model. The data could then be used in a geometry shader, where each coordinate is transformed into a cube. This avoids the copying to the CPU and requires less data to store the voxelization, since only occupied voxels are stored. It can, however, render the same voxel multiple times if the voxelization does not keep track of which voxels are already occupied.

Therefore, rendering the scene using raymarching was both deemed easier and a potentially faster method. The data was not transferred between the CPU and GPU and only the visible voxels were rendered.

To perform the raymarching, it used the RLV algorithm to traverse the voxels. In this case however, the line was not terminated when it reached the endpoint of the line segment, as it does not exist. Instead it terminated when the line intersected with an existing voxel or when it had exited the voxel grid.

4.2 Evaluation

When performing the comparisons of the different line algorithms, several variable changes were considered in the experiments.

Firstly, multiple models were tested on, with ranging levels of detail. The models that were used included the Blender monkey (also called Suzanne) [24], the Stanford bunny [25] and the Stanford dragon [25]. These can be seen in Figure 4.5. The triangle count of each model was 3936, 69451 and 871414 respectively. Suzanne was also subdivided once in Blender [26], in order to increase the triangle count to 3936.

Secondly, the resolution of the voxel grid was varied for each model. The resolution varied between 128-2048, incremented by powers of 2. Due to GPU memory limits, resolutions greater than 2048 were not possible.




Figure 4.5: Models used for evaluating the voxelizations

4.2.1 Performance Analysis

To profile the performance of the algorithms, CUDA events were used [27]. To use the CUDA events, two timestamp events were created. One of them was initialized before the kernel ran using the CUDA function `cuEventRecord`. Then after the kernel ran, another event was initialized. These timestamps were entirely handled by the GPU, meaning the CPU would not wait for the timestamps to be executed. The function `cuEventSynchronize`, was used in order to wait for the events to be executed. The time between two timestamps was then determined by the CUDA call `cuEventElapsedTime`.

4.2.2 Error Analysis

The error was calculated for each of the resolutions, models and combinations of line voxelizations. Meaning RLV was compared to both ILV and Bresenham, but ILV was also compared to Bresenham.

The comparison was done by first voxelizing using the first algorithm with the voxel value of 1. Then the voxelization ran again with the other algorithm (with the same 3D texture). This time the voxel value was determined depending on the current voxel value in that position. If the value at that position was 1 or 2, the value 2 was written, otherwise the value 3 was written. The result was a voxelization where the intersection had value 2, the voxels only in the first algorithm had value 1 and the voxels only in the second had value 3.

To calculate the error, a simple CUDA kernel was created which iterated through the whole voxel space and summed up the amount of each of the values. With these values the error could be calculated using the formulas in Section 3.3.

5

Results

Here the results of the thesis will be presented. This includes figures of the voxelization, graphs of the performance and data of the error analysis.

5.1 Voxelization

The result of the floating-point voxelization can be seen in Figure 5.1. The same voxelization of the integer versions is shown in Figure 5.2 and Figure 5.3. The figures show voxelizations between 16-512 resolution, but resolutions up to 2048 were possible. However, due to the aliasing in the rendering, showing the result of such high resolution would be meaningless.

All the models in the voxelization were scaled to fit the voxel grid perfectly, without changing the aspect ratio of them. More figures of the voxelizations can be found in Appendix A.4.

5.2 Performance Analysis

The performance of each of the models can be seen in Figure 5.4. The same data is plotted for each of the algorithms in Figure 5.5. Each data point is the average time to voxelize the models in a total of 100 iterations. The raw performance data can be found in Appendix A.5.

One unsolved problem about the performance data is that the timings improved after iterating the voxelization multiple times. As an example, Figure 5.6 shows how the timings varied after a certain amount of iterations. The time for a single voxelization decreased by around 50%. This will further be discussed in the next chapter. Since the data seem to converge to a certain time, the average

of the last 100 out of the 1000 iterations were used as the final measurement in the previously mentioned graphs.




Figure 5.1: Floating-point voxelization using RLV of the Stanford bunny at 16, 32, 64, 128, 256 and 512 resolution




Figure 5.2: Integer voxelization using ILV of the Stanford bunny at 16, 32, 64, 128, 256 and 512 resolution




Figure 5.3: Integer voxelization using Bresenham of the Stanford bunny at 16, 32, 64, 128, 256 and 512 resolution




Figure 5.4: Performance of the voxelization for a model with different algorithms and resolutions




Figure 5.5: Performance of the voxelization for the algorithms with varying models and resolutions




Figure 5.6: Performance data of each iteration for a total of 1000 iterations. This was performed using the Stanford bunny with 512 voxel resolution. Similar data is given for different models and resolutions.

5.3 Error Analysis

An example of the difference between RLV and ILV can be seen in Figure 5.7. It shows a comparison where voxels are colored depending on if they exist in both voxelizations or only one of them. More comparisons can be seen in Appendix A.6, where all models and resolutions between 16-512 are shown.

The Jaccard distance between the different versions of the voxelizations can be seen in Tabel 5.1.




Figure 5.7: Difference between floating-point and integer voxelization. The left-most figure uses RLV, the middle uses ILV and the right-most is the union between them. The yellow voxels are in both versions. The red voxels are only in the RLV version. The blue voxels are only in the ILV version.

Model	Algorithm	128	256	512	1024	2048
Monkey	RLV/ILV	21.63%	22.10%	22.51%	22.98%	23.53%
	RLV/Bre	21.89%	22.30%	23.04%	23.67%	24.16%
	ILV/Bre	4.72%	5.62%	5.58%	5.35%	5.21%
Bunny	RLV/ILV	19.65%	20.62%	22.12%	22.59%	22.8%
	RLV/Bre	19.66%	21.21%	22.30%	22.94%	23.21%
	ILV/Bre	0.07%	4.57%	4.77%	5.59%	5.23%
Dragon	RLV/ILV	9.42%	15.65%	21.28%	21.87%	22.75%
	RLV/Bre	9.42%	15.67%	21.40%	22.52%	23.29%
	ILV/Bre	0.00%	0.05%	1.09%	3.84%	5.22%

Table 5.1: The Jaccard distance between the different algorithms, with varying models and resolution. Bre in the table is short for Bresenham. All errors are rounded to the nearest hundredth.

6

Discussion

This chapter discusses the results from the method. This includes explaining the given data and how the data changes depending on models and resolution. It will also give a critical view of the method and explain why certain methods were used. Finally, a discussion will be had about future improvements and research which can be done surrounding the thesis.

6.1 Results

This section discusses the results shown in the previous chapter. It serves the purpose of explaining and evaluating the data gathered from the method.

6.1.1 Performance of Models

From Figure 5.5 in the results, it can be seen that the voxelization generally takes longer with the increase in triangle count. This makes sense, as having more triangles to process should increase the amount of time to voxelize. The timings of the models do however seem to converge towards the same value as the resolution increases. This is likely due to the fact that even though there are less triangles in the monkey compared to the dragon, the monkey needs to generate more voxels per triangle. Meaning the GPU cannot utilize its parallelism when less triangles are being voxelized. This can be seen in Figure 5.5 for RLV, where even though the dragon has ten times more triangles, it outperforms the bunny at higher resolutions.

6.1.2 Performance Between RLV and ILV

Looking back at article of the optimal scanline [9], the authors claimed ILV was around 3% faster than RLV. Again, the data was not published in the article. This is however not consistent with the results of this thesis, where ILV was worse by between 20-250%. There could be several reasons for this.

First of, the article never states anything about the edge-cases in the integer version presented in Section 4.1.3. Fixing these edge-cases caused some overhead, which they might not have had.

Secondly, much of the focus of the article was around the integer version and few details were explained about the floating-point version. This meant a lot of improvisation had to be done when implementing the floating-point version. Therefore, it is possible it was implemented in a more efficient way.

Thirdly, there is of course a possibility that the authors skipped some optimization steps in the article. This could be because they had a budget on the amount of pages they could write, or the full optimization steps were deemed too complicated for the article.

The results of my implementation are however reasonable considering the different complexities of finding the scanline endpoints. In the floating-point version, all that had to be done is increasing the scanline length and calculating a reverse projection of the scanline direction. The integer version required iterating through multiple voxels to find the endpoint. As such, it would make sense that the floating-point version required fewer calculations to find the endpoints.

6.1.3 Performance Between ILV and Bresenham

Since ILV and Bresenham use the same integer version, the performance difference is solely based on the line voxelization algorithm. The original hypothesis was that Bresenham would improve performance, but looking at the data, this is not the case.

The big reason for this was the last problem described in Section 4.1.4, that Bresenham could generate three voxels per iteration. This meant that additional overhead was needed in order to get the next voxel in the Bresenham voxelization. In ILV the same operations are performed no matter which direction the next voxel is in, but in Bresenham different calculations are done at different stages. This resulted in Bresenham causing more branch divergence on the GPU, which caused a major slow down.

6.1.4 Performance Data

After analyzing the performance data, it was noted that the timings changed drastically when iterating the voxelization many times. The reason for this is still unknown, but it is somehow caused by AWS. This was discovered after testing the voxelization on normal desktop computers, which had none of these issues. There are two hypotheses as to why this could happen.

One reason could be that the program ran over a network. As described in Section 2.2, to run the voxelization, ThinLinc together with VirtualGL was used.

This could potentially cause extra overhead which brings down the performance. It does however not explain why the CUDA performance is suffering from this, as VirtualGL only enables OpenGL 3D acceleration. However as the implementation makes use of OpenGL textures, VirtualGL could still be the culprit.

The other hypothesis is that AWS allocates more resources as it is needed. Meaning the program does not run at full power until it has run for a certain amount of time. This could explain the distinct performance boosts seen in Figure 5.6 in the results.

Again, there is no conclusive answer as to why the performance increases, the reasons above are just speculations as to why.

6.1.5 Error Analysis

Tabel 5.1 in the results shows all the error data between the different voxelization algorithms. The error between RLV and ILV grew as the resolution of the voxelization increased. The same was true between RLV and Bresenham. It seems like the error converges around 25%, which is a considerable difference to the 2.5% presented by Zhang et al. [9]. Again, the error presented in the article is a relative error of the voxel count between the two voxelizations, which is different to the Jaccard distance.

An interesting data point is the error for the Stanford dragon. In the lower resolutions the error was a lot less than the error for the other models. This is likely due to the fact that the triangle count was much higher. The result is that a lot of triangles were fully inside a voxel, meaning there could not be an error for those triangles. For the triangles which were not fully inside a voxel, they generally spanned a total of 3-4 voxels and were therefore also prone to less error. This is also consistent with the error increasing as the resolution increased.

Looking at Tabel 5.1, it can be seen that there was an error of up to 5% between ILV and Bresenham. As they both used the same underlying integer algorithm, it would seem odd that there was an error between them. The error occurred when the line touched more than two voxels. That is, the line did not pass through the face of the voxels. This is shown in Figure 6.1. The ambiguity of which voxel should be chosen caused the two versions to generate slightly different voxelizations. A potential fix for this would be to voxelize all the voxels touching the line, but this would most likely cause more edge-cases when finding the next scanline endpoint.




Figure 6.1: Line touching the corner of four voxels. This causes ambiguity in whether the dark gray or the white voxel should be voxelized.

6.2 Method

This section serves the purpose of motivating and critically analysing the methods used in the thesis.

6.2.1 Performance Analysis

To run the performance analysis, CUDA events were the choice for this thesis. These events only measured the execution time of the kernel and did not include the time it took to launch it. An alternative to this would be to use CPU timers. This would give a more real depiction of how long the voxelization would take in a real environment, as it would include the launch time of the kernel. However, this thesis is set out to measure the performance of the voxelization and including the launch time of the kernel was deemed unnecessary.

6.2.2 Error Analysis

There is also something to be said about how the error was measured. As the Jaccard distance only calculates the amount of mismatched voxels, it is not a measurement of how well the general shape of the voxelizations matches. As such, it should not be used for comparing dissimilarity of voxelizations between different models. It is more of a metric to compare already similar data sets. The results of using this method could therefore be misleading when saying the voxelization has an error of 20-25%. Such great error might leave the reader to believe the integer version is unrecognizable to the floating-point version. This is however not the case as the general shape is preserved very well, and it is hard to visually see the error without a comparison figure.

6.3 Future Work

As always with a project like this there is still room for improvements. Some of those improvements are described in the following paragraphs.

Firstly, modify the Bresenham algorithm to perform better when retrieving the next voxel in the line. The problem is that it performs different calculations depending on which axis it is currently evaluating. One solution could be to add an additional error for the dominant axis. This would then result in all axes performing the same calculations. Then to get the next voxel iterate the different axes until a voxel is found. To determine which axis to evaluate, a variable can be used to index the axes.

Secondly, analyze the error between the algorithms when performing a supercover line voxelization. A supercover avoids the ambiguity of which voxel to voxelize when the line touches multiple voxels. This is done by voxelizing all the voxels instead of choosing one of them. Doing this could fix the problem with the integer versions generating different voxelizations due to the ambiguity.

7

Conclusion

This chapter answers the research questions by summarizing the results and discussions. It also gives a conclusive answer as to which algorithm one should use when implementing an optimal scanline algorithm.

7.1 Research Questions

The research questions presented in the introduction are reiterated and answered in the following paragraphs.

1. Which line voxelization algorithm performs best for the optimal scanline?

For all models and voxel grid resolution, the floating-point version, RLV, performed better. ILV performed anywhere between 20-250% worse compared to RLV. Bresenham performed anywhere between 20-500% worse compared to RLV. Interestingly, ILV performed better than Bresenham in every case but two. Although, the difference was not as dramatic as between RLV and ILV, except for a few outliers which performed way worse.

The reason for the big difference between the floating-point and integer versions was due to the complexity of finding the scanline endpoints. For RLV, it was simply a case of increasing the length of the scanline direction and performing a reverse projection of the direction to the triangle's edges. For ILV and Bresenham, finding the endpoints required iterating through the edge voxelizations and performing a boundary test based on the previous scanline and the current voxel.

The reason for ILV being faster than Bresenham had less to do with Bresenham being worse in general. It had more to do with Bresenham not being able to step the line voxelization by one voxel each iteration. This in turn required extra overhead to overcome.

2. How great is the approximation error of the integer versions of the optimal scanline?

The error for the integer versions of the optimal scanline was determined to be around 20-25%. This was based on the Jaccard distance between RLV and ILV/Bresenham. One outlier of the data was the dragon model, with 871414 triangles, where the error started at around 9% for a voxel grid resolution of 128. It however gradually grew as the resolution increased. The reason was likely due to the triangles of the model being mostly within a single or a few voxels. Making it less likely to cause an error.

There also turned out to be a minor error (~5%) between ILV and Bresenham, even though they both used the same underlying algorithm. This was caused due to cases where the scanline passes through multiple voxels, such that the choice of voxel was ambiguous. An example of this was shown in Figure 6.1.

7.2 Choice of Algorithm

From the results given there is only one obvious choice as to which algorithm one should use. This is of course the floating-point version using RLV. The algorithm had the best performance in terms of runtime. It is also closer to the ground truth, as it only voxelizes voxels touching the triangle. Due to approximation errors, this is not the case for the integer versions. Given those results, there should be no reason not to choose RLV.

Appendix

A

Appendix

A.1 3D Bresenham Algorithm

```
 $e_y \leftarrow 2\Delta Y - \Delta X$ 
 $e_z \leftarrow 2\Delta Z - \Delta X$ 
 $x \leftarrow X_0, y \leftarrow Y_0, z \leftarrow Z_0$ 
while  $x \leq X_1$  do
    |  $voxel \leftarrow (x, y, z)$ 
    | SetVoxel( $voxel$ )
    |  $x \leftarrow x + 1$ 
    | if  $e_y \geq 0$  then
    |   |  $y \leftarrow y + 1$ 
    |   |  $e_y \leftarrow e_y + 2(\Delta Y - \Delta X)$ 
    | else
    |   |  $e_y \leftarrow e_y + 2\Delta Y$ 
    | end
    | if  $e_z \geq 0$  then
    |   |  $z \leftarrow z + 1$ 
    |   |  $e_z \leftarrow e_z + 2(\Delta Z - \Delta X)$ 
    | else
    |   |  $e_z \leftarrow e_z + 2\Delta Z$ 
    | end
| end
```

Algorithm 1: The Bresenham algorithm for calculating a straight line in 3D

A.2 6-Connected Bresenham Algorithm Modification

```

if  $e_y \geq 0$  then
    if  $e_y \geq \Delta Y$  then
        | voxel  $\leftarrow$  (voxel.x, voxel.y + 1, voxel.z)
    else
        | voxel  $\leftarrow$  (x, y, z)
    end
    SetVoxel(voxel)
    y  $\leftarrow$  y + 1
     $e_y \leftarrow e_y + 2(\Delta Y - \Delta X)$ 
else
    |  $e_y \leftarrow e_y + 2\Delta Y$ 
end

```

Algorithm 2: Modification for the y-axis of the Bresenham algorithm to voxelize a 6-connected line. The same modification is done for the z-axis.

A.3 CUDA-OpenGL Interoperability

```
1 // ===== DEVICE =====
2
3 surface<void, 3> voxelGrid;
4
5 __device__
6 void WriteToTexture(int3 voxel, unsigned char color)
7 {
8     surf3Dwrite(color, voxelGrid, voxel.x, voxel.y, voxel.z);
9 }
10
11 // ===== HOST =====
12
13 // Create OpenGL texture
14 glGenTextures(1, &glTex);
15 glBindTexture(GL_TEXTURE_3D, glTex);
16 glTexParameteri(GL_TEXTURE_3D, GL_TEXTURE_MIN_FILTER, GL_NEAREST);
17 glTexParameteri(GL_TEXTURE_3D, GL_TEXTURE_MAG_FILTER, GL_NEAREST);
18 glTexImage3D(GL_TEXTURE_3D, 0, GL_R8, size, size, size, 0, GL_RED,
19             GL_FLOAT, nullptr);
20
21 // Register the OpenGL texture as a CUDA texture
22 cuGraphicsGLRegisterImage(&cudaTex, glTex, GL_TEXTURE_3D,
23                           CU_GRAPHICS_REGISTER_FLAGS_SURFACE_LDST);
24
25 // Bind the CUDA texture to a CUDA array
26 cuGraphicsMapResources(1, &cudaTex, 0);
27 cuGraphicsSubResourceGetMappedArray(&cudaArray, cudaTex, 0, 0);
28 cuGraphicsUnmapResources(1, &cudaTex, 0);
29
30 // Link the CUDA array to a CUDA surface
31 cuModuleGetSurfRef(&cudaSurfRef, module, "voxelGrid");
32 cuSurfRefSetArray(cudaSurfRef, cudaArray, 0);
```

A.4 Voxelizations




Figure A.1: Floating-point voxelization using RLV of the Blender monkey at 16, 32, 64, 128, 256 and 512 resolution




Figure A.2: Integer voxelization using ILV of the Blender monkey at 16, 32, 64, 128, 256 and 512 resolution




Figure A.3: Integer voxelization using Bresenham of the Blender monkey at 16, 32, 64, 128, 256 and 512 resolution




Figure A.4: Floating-point voxelization using RLV of the Stanford dragon at 16, 32, 64, 128, 256 and 512 resolution




Figure A.5: Integer voxelization using ILV of the Stanford dragon at 16, 32, 64, 128, 256 and 512 resolution




Figure A.6: Integer voxelization using Bresenham of the Stanford dragon at 16, 32, 64, 128, 256 and 512 resolution

A.5 Performance Data

Model	Algorithm	128	256	512	1024	2048
Monkey	RLV	0.168047	0.360608	1.17848	4.18106	16.5649
	ILV	0.213802	0.576965	1.84135	5.96018	21.9102
	Bre	0.24961	0.636628	1.86863	5.72366	19.9131
Bunny	RLV	0.193737	0.369392	1.76889	6.86508	21.5765
	ILV	0.602098	1.22562	3.29431	9.25156	27.0328
	Bre	0.988574	1.92251	4.5347	11.6089	31.9672
Dragon	RLV	1.22597	1.50882	2.32364	4.89038	18.1659
	ILV	3.86642	5.36262	8.12676	16.3983	42.3611
	Bre	6.71438	9.32611	13.8366	26.5671	60.6166

Table A.1: The Raw performance data of the different algorithms, with varying models and resolution. Bre in the table is short for Bresenham. Timings are in milliseconds.

A.6 Voxelization Error

In the following figures the yellow voxels are voxels in both algorithms. The red voxels are only in the algorithm mentioned first and the blue voxels are only in the algorithm mentioned last. For example, in Figure A.7, RLV is mentioned first and therefore correspond with red voxels, while ILV is mentioned last and are blue voxels.




Figure A.7: Difference between RLV and ILV for the Blender monkey at 16, 32, 64, 128, 256 and 512 resolution




Figure A.8: Difference between RLV and Bresenham for the Blender monkey at 16, 32, 64, 128, 256 and 512 resolution




Figure A.9: Difference between ILV and Bresenham for the Blender monkey at 16, 32, 64, 128, 256 and 512 resolution




Figure A.10: Difference between RLV and ILV for the Stanford bunny at 16, 32, 64, 128, 256 and 512 resolution




Figure A.11: Difference between RLV and Bresenham for the Stanford bunny at 16, 32, 64, 128, 256 and 512 resolution




Figure A.12: Difference between ILV and Bresenham for the Stanford bunny at 16, 32, 64, 128, 256 and 512 resolution




Figure A.13: Difference between RLV and ILV for the Stanford dragon at 16, 32, 64, 128, 256 and 512 resolution




Figure A.14: Difference between RLV and Bresenham for the Stanford dragon at 16, 32, 64, 128, 256 and 512 resolution




Figure A.15: Difference between ILV and Bresenham for the Stanford dragon at 16, 32, 64, 128, 256 and 512 resolution

Bibliography

- [1] J. Amanatides and A. A. Woo. A fast voxel traversal algorithm for ray tracing. In *Eurographics*, volume 87, pages 3–10, 1987. Cited on pages iii, 2, and 10.
- [2] C. Crassin, F. Neyret, M. Sainz, S. Green, and E. Eisemann. Interactive indirect illumination using voxel cone tracing. *Comput. Graph. Forum*, 30:1921–1930, 09 2011. doi: 10.1111/j.1467-8659.2011.02063.x. Cited on page 1.
- [3] William E Lorensen and Harvey E Cline. Marching cubes: A high resolution 3d surface construction algorithm. *ACM siggraph computer graphics*, 21(4):163–169, 1987. Cited on page 1.
- [4] J. M. L. Guevara and D. R. Albrecht. From pixel to voxel in the clinical diagnosis. In *2016 IEEE 36th Central American and Panama Convention (CONCAPAN XXXVI)*, pages 1–5, 11 2016. doi: 10.1109/CONCAPAN.2016.7942370. Cited on page 1.
- [5] D. Zhang and M. M. F. Yuen. Collision detection for clothed human animation. In *Proceedings the Eighth Pacific Conference on Computer Graphics and Applications*, pages 328–337, 10 2000. doi: 10.1109/PCCGA.2000.883956. Cited on page 1.
- [6] Michael Schwarz and Hans-Peter Seidel. Fast parallel surface and solid voxelization on gpus. *ACM transactions on graphics (TOG)*, 29(6):1–10, 2010. Cited on page 1.
- [7] C. Crassin and S. Green. *Octree-Based Sparse Voxelization Using the GPU Hardware Rasterizer*, pages 303–320. 07 2012. ISBN 978-1-4398-9376-0. doi: 10.1201/b12288-26. Cited on pages 1 and 14.
- [8] Y. Fei, B. Wang, and J. Chen. Point-tessellated voxelization. In *Proceedings of Graphics Interface 2012, GI ’12*, page 9–18, CAN, 2012. Canadian Information Processing Society. ISBN 9781450314206. Not cited.
- [9] Y. Zhang, S. Garcia, W. Xu, T. Shao, and Y. Yang. Efficient voxelization using projected optimal scanline. *Graphical Models*, 100:61 – 70, 2018. ISSN 1524-

0703. doi: <https://doi.org/10.1016/j.gmod.2017.06.004>. Cited on pages 1, 2, 11, 12, 14, 15, 32, and 33.
- [10] X. Liu and K. Cheng. Three-dimensional extension of bresenham's algorithm and its application in straight-line interpolation. *Proceedings Of The Institution of Mechanical Engineers Part B-Journal of Engineering Manufacture - PROC INST MECH ENG B-J ENG MA*, 216:459–463, 03 2002. doi: 10.1243/0954405021519979. Cited on pages 2, 11, and 12.
- [11] The Khronos Group Inc. Opengl overview, 2020. URL <https://www.opengl.org/about/>. Visited on 2020-02-11. Cited on page 6.
- [12] Apple Inc. About opengl for os x, 2018. URL https://developer.apple.com/library/archive/documentation/GraphicsImaging/Conceptual/OpenGL-MacProgGuide/opengl_intro/opengl_intro.html. Visited on 2020-02-10. Cited on page 6.
- [13] Apple Inc. About opengl es, 2018. URL https://developer.apple.com/library/archive/documentation/3DDrawing/Conceptual/OpenGLES_ProgrammingGuide/Introduction/Introduction.html. Visited on 2020-02-10. Cited on page 6.
- [14] NVIDIA Corporation. Nvcc, 2019. URL <https://docs.nvidia.com/cuda/cuda-compiler-driver-nvcc/index.html>. Visited on 2020-02-19. Cited on page 6.
- [15] Amazon. Amazon web services, 2020. URL <https://aws.amazon.com/>. Visited on 2020-05-11. Cited on page 7.
- [16] Cendio. Thinlinc - a remote desktop server from cendio, 2020. URL <https://www.cendio.com/thinlinc/what-is-thinlinc>. Visited on 2020-04-29. Cited on page 7.
- [17] Virtual GL. The virtualgl project, 2018. URL <https://www.virtualgl.org/>. Visited on 2020-04-29. Cited on page 7.
- [18] id Software. Wolfenstein 3d, 2012. URL <https://github.com/id-Software/wolf3d>. Visited on 2020-02-19. Cited on page 8.
- [19] J. E. Bresenham. Algorithm for computer control of a digital plotter. *IBM Systems Journal*, 4(1):25–30, 1965. ISSN 0018-8670. doi: 10.1147/sj.41.0025. Cited on page 11.
- [20] A. Karabassi, G. Papaioannou, and T. Theoharis. A fast depth-buffer-based voxelization algorithm. *Journal of Graphics Tools*, 4:5–10, 01 1999. doi: 10.1080/10867651.1999.10487510. Cited on pages 1 and 15.
- [21] P. Jaccard. Distribution de la flore alpine dans le bassin des dranses et dans quelques régions voisines. *Bull Soc Vaudoise Sci Nat*, 37:241–272, 1901. Cited on page 16.

- [22] Tim Håkansson. Master thesis, 2020. URL <https://github.com/thraix/MasterThesis>. Visited on 2020-06-06. Cited on page 17.
- [23] D. Sandler. gl_cuda_interop_pingpong_st, 2015. URL https://github.com/nvpro-samples/gl_cuda_interop_pingpong_st. Visited on 2020-03-10. Cited on page 17.
- [24] Blender Foundation. Primitives, 2020. URL <https://docs.blender.org/manual/en/latest/modeling/meshes/primitives.html>. Visited on 2020-06-02. Cited on page 22.
- [25] Stanford University. The stanford 3d scanning repository, 2014. URL <https://graphics.stanford.edu/data/3Dscanrep/>. Visited on 2020-06-02. Cited on page 22.
- [26] Blender Foundation. Blender 2.82a, 2020. URL <https://www.blender.org>. Visited on 2020-06-02. Cited on page 22.
- [27] M. Harris. How to implement performance metrics in cuda c/c++, 2012. URL <https://devblogs.nvidia.com/how-implement-performance-metrics-cuda-cc/>. Visited on 2020-03-24. Cited on page 22.